


TaskBench: Model Quality Across 21 Production Tasks



46 models, 21 tasks. V2 judge (GPT-4o, correctness-focused). Label colors indicate price tier: red = Premium (\$5+/M), blue = Standard (\$0.50 - \$4.99), green = Economy (\$0.05 - \$0.49), orange = Micro (< \$0.05).