


Judge V2 vs HumanEval pass@1



Pearson $r = 0.883$, $p < 0.001$ (significant). $n = 48$ models (includes incomplete models). Dashed line: linear fit.